





# ETC1010: Data Modelling and Computing

## Lecture 3: Wrangling your data





Di Cook ([dicook@monash.edu](mailto:dicook@monash.edu), @visnut)

Week 3

# Overview




-  Data structures, variable types
-  Wrangling verbs: filter, arrange, select, mutate, summarise
-  Making joins
-  Working with dates

# Data structures



-  Data frames
-  Matrices/vectors
-  Lists
-  Tibbles

# data.frame/tibble vs matrices/vectors





Mostly what we have seen so far are data.frame or tibble

-  Rectangular format
-  Each column might be a different type of variable
-  Each column is same length

Matrices/vectors

-  Original data format
-  All columns have numerical values

# Lists

-  Not necessarily rectangular
-  Anything can be packed into a list
-  We will see more of these when we fit models, because model summaries are typically shared as a list
-  Trickier to work with

# How do you know what you have?

```
genes <- read_csv("data/genes.csv")
genes
# A tibble: 3 x 12
  id `WI-6.R1` `WI-6.R2` `WI-6.R4` `WM-6.R1` `WM-6.R2` `WI-12.R1`
  <chr>      <dbl>      <dbl>      <dbl>      <dbl>      <dbl>      <dbl>
1 Gene 1    2.182424  2.2042219  4.195636  2.6273345  5.063641  4.540002
2 Gene 2    1.464224  0.5854472  1.859238  0.5152242  2.882808  1.364037
3 Gene 3    2.031792  0.8701078  3.281983  0.5330452  4.627315  2.182192
# ... with 5 more variables: `WI-12.R2` <dbl>, `WI-12.R4` <dbl>,
#   `WM-12.R1` <dbl>, `WM-12.R2` <dbl>, `WM-12.R4` <dbl>
is_tibble(genes)
[1] TRUE
is.data.frame(genes)
[1] TRUE
is.matrix(genes)
[1] FALSE
is.list(genes)
[1] TRUE
```

```

melbtemp <- read.fwf("data/ASN00086282.dly",
  c(11, 4, 2, 4, rep(c(5, 1, 1, 1), 31)), fill=T)
head(melbtemp[,c(1,2,3,4,seq(5,100,4))])

```


		V1	V2	V3	V4	V5	V9	V13	V17	V21	V25	V29	V33	V37	V41	V45	V49
1	ASN00086282	1970	7	TMAX	141	124	113	123	148	149	139	153	123	108	119	112	
2	ASN00086282	1970	7	TMIN	80	63	36	57	69	47	84	78	49	42	48	56	
3	ASN00086282	1970	7	PRCP	3	30	0	0	36	3	0	0	10	23	3	0	
4	ASN00086282	1970	8	TMAX	145	128	150	122	109	112	116	142	166	127	117	127	
5	ASN00086282	1970	8	TMIN	50	61	75	67	41	51	48	-7	56	62	47	33	
6	ASN00086282	1970	8	PRCP	0	66	0	53	13	3	8	0	0	0	3	5	


```

  V53 V57 V61 V65 V69 V73 V77 V81 V85 V89 V93 V97
1 126 112 115 133 134 126 104 143 141 134 117 142
2  51  36  44  39  40  58  15  33  51  74  39  66
3   5   0   0   0   0   0   8   0  18   0   0   0
4 159 143 114  65 113 125 129 147 161 168 178 161
5  67  84  11  41  18  50  22  28  74  94  73  88
6   0   0  64   3  99  36   8   0   0   0   8  36
is_tibble(melbtemp)
[1] FALSE
is.data.frame(melbtemp)
[1] TRUE
is.matrix(melbtemp)
[1] FALSE
is.list(melbtemp)
[1] TRUE


```

# Variable types


 `int` integer.

 `dbl` doubles, or real numbers

 `num` numeric, includes integer or double

 `chr` characters, or strings

 `lgl` or `logi` logical, true or false

 `fctr` or `Factor` fancy character, contains additional information about levels and labels

 `date` and `dtm` time and date structures



# Converting formats

```
melbtemp <- melbtemp[,c(1,2,3,4,seq(5,128,4))]  
colnames(melbtemp) <- c("id", "year", "month", "var", paste0("V",1:31))  
melbtemp <- melbtemp %>%  
  gather(day, value, V1:V31) %>%  
  mutate(day = sub("V", "", day)) %>%  
  mutate(value=ifelse(value== -9999, NA, value)) %>%  
  filter(var %in% c("PRCP", "TMAX", "TMIN")) %>%  
  spread(var, value) %>%  
  mutate(PRCP=PRCP/10, TMAX=TMAX/10, TMIN=TMIN/10) %>%  
  as_tibble()  
melbtemp  
# A tibble: 16,399 x 7  
      id year month   day PRCP  TMAX  TMIN  
  <fctr> <int> <int> <chr> <dbl> <dbl> <dbl>  
1 ASN00086282 1970     7     1   0.3  14.1   8.0  
2 ASN00086282 1970     7    10   2.3  10.8   4.2  
3 ASN00086282 1970     7    11   0.3  11.9   4.8  
4 ASN00086282 1970     7    12   0.0  11.2   5.6  
5 ASN00086282 1970     7    13   0.5  12.6   5.1  
6 ASN00086282 1970     7    14   0.0  11.2   3.6  
7 ASN00086282 1970     7    15   0.0  11.5   4.4  
8 ASN00086282 1970     7    16   0.0  13.3   3.9  
9 ASN00086282 1970     7    17   0.0  13.4   4.0  
10 ASN00086282 1970     7    18   0.0  12.6   5.8  
# ... with 16,389 more rows
```

# Filter


Pick observations by their values. For example,


```
filter(var %in% c("PRCP", "TMAX", "TMIN"))
```


took the column named `var` and keeps rows that have one of three values `PRCP`, `TMAX`, `TMIN`. All other rows are removed.


# Logical operators

## Comparisons

  $<$ ,  $<=$ : less than, less than or equal

  $>$ ,  $>=$ : greater than, greater than or equal

  $==$ ,  $!=$ : equal to, not equal to

 Cautions: computers have a hard time with equals, `near()` is a function that can be useful for real numbers, will check within a small neighbourhood of a value

## Logicals

  $\&$  and, both checks have to be true

  $|$  or, either check is true

 `%in%` checks the elements of a collection, multiple or's

## Missings: `is.na()`

# Arrange

Orders a data frame or tibble by the values in one column

```
library(lubridate)
melbtemp %>%
  mutate(date=ymd(paste(year, month, day, sep="-"))) %>%
  arrange(date)
# A tibble: 16,399 x 8
```

	id	year	month	day	PRCP	TMAX	TMIN	date
	<fctr>	<int>	<int>	<chr>	<dbl>	<dbl>	<dbl>	<date>
1	ASN00086282	1970	7	1	0.3	14.1	8.0	1970-07-01
2	ASN00086282	1970	7	2	3.0	12.4	6.3	1970-07-02
3	ASN00086282	1970	7	3	0.0	11.3	3.6	1970-07-03
4	ASN00086282	1970	7	4	0.0	12.3	5.7	1970-07-04
5	ASN00086282	1970	7	5	3.6	14.8	6.9	1970-07-05
6	ASN00086282	1970	7	6	0.3	14.9	4.7	1970-07-06
7	ASN00086282	1970	7	7	0.0	13.9	8.4	1970-07-07
8	ASN00086282	1970	7	8	0.0	15.3	7.8	1970-07-08
9	ASN00086282	1970	7	9	1.0	12.3	4.9	1970-07-09
10	ASN00086282	1970	7	10	2.3	10.8	4.2	1970-07-10

```
# ... with 16,389 more rows
```

# Select

Choose some of the **variables**.

```
airport <- read_csv("data/airports.csv")
airport
# A tibble: 13,094 x 29
  AIRPORT_SEQ_ID AIRPORT_ID AIRPORT      DISPLAY_AIRPORT_NAME
      <int>      <int>    <chr>          <chr>
1     1000101     10001      01A      Afognak Lake Airport
2     1000301     10003      03A      Bear Creek Mining Strip
3     1000401     10004      04A      Lik Mining Camp
4     1000501     10005      05A      Little Squaw Airport
5     1000601     10006      06A      Kizhuyak Bay
6     1000701     10007      07A      Klawock Seaplane Base
7     1000801     10008      08A      Elizabeth Island Airport
8     1000901     10009      09A      Augustin Island
9     1001001     10010      1B1      Columbia County
10    1001002     10010      1B1      Columbia County
# ... with 13,084 more rows, and 25 more variables:
#   DISPLAY_AIRPORT_CITY_NAME_FULL <chr>, AIRPORT_WAC <int>,
#   AIRPORT_COUNTRY_NAME <chr>, AIRPORT_COUNTRY_CODE_ISO <chr>,
#   AIRPORT_STATE_NAME <chr>, AIRPORT_STATE_CODE <chr>,
#   AIRPORT_STATE_FIPS <chr>, CITY_MARKET_ID <int>,
#   DISPLAY_CITY_MARKET_NAME_FULL <chr>, CITY_MARKET_WAC <int>,
#   LAT_DEGREES <int>, LAT_HEMISPHERE <chr>, LAT_MINUTES <int>,
#   LAT_SECONDS <int>, LATITUDE <dbl>, LON_DEGREES <int>,
#   LON_HEMISPHERE <chr>, LON_MINUTES <int>, LON_SECONDS <int>,
```


```
airport %>%
  select(AIRPORT, LATITUDE, LONGITUDE, AIRPORT_IS_LATEST, DISPLAY_AIRPORT_NAME)
# A tibble: 13,094 x 5
  AIRPORT LATITUDE LONGITUDE AIRPORT_IS_LATEST DISPLAY_AIRPORT_NAME
  <chr>     <dbl>     <dbl>         <int>         <chr>
1    01A  58.10944 -152.90667           1 Afognak Lake Airport
2    03A  65.54806 -161.07167           1 Bear Creek Mining Strip
3    04A  68.08333 -163.16667           1 Lik Mining Camp
4    05A  67.57000 -148.18389           1 Little Squaw Airport
5    06A  57.74528 -152.88278           1 Kizhuyak Bay
6    07A  55.55472 -133.10167           1 Klawock Seaplane Base
7    08A  59.15694 -151.82917           1 Elizabeth Island Airport
8    09A  59.36278 -153.43056           1 Augustin Island
9    1B1  42.28889 -73.71028            0 Columbia County
10   1B1  42.29139 -73.71028           1 Columbia County
# ... with 13,084 more rows
```

# Mutate

Create new, or transform existing, variables, e.g. for the Melbourne temperature data, we put the temperature and precipitation values into Celsius and mm, byt dividing by 10.

```
mutate(PRCP=PRCP/10, TMAX=TMAX/10, TMIN=TMIN/10)
```


# Arithmetic

 Typical calculations:  $+$ ,  $-$ ,  $*$ ,  $/$ ;  $^$  (power)


 Modular:

  $\%/\%$  integer division

  $\%\%$  remainder

 Transformations:  $\log()$ ,  $\log_{10}()$ ,  $\sqrt{\phantom{x}}$

 Temporal:  $\text{lead}()$ ,  $\text{lag}()$  create variables which are temporal lags of existing


 Stats:  $\text{rank}()$ ,  $\text{cumsum}()$




# Summarise

Calculate a quantity on a column, producing a single number, e.g. for the audio data:

```
summarise(m = mean(value), s = sd(value),  
          mx = max(value), mn = min(value))
```

 Stats: mean(), median(), sd(), min(), max(), sum()

 Rounding: floor(), ceiling(), trunc(), round(), signif()

# Grouping and ungrouping

Summarise is most commonly called on subgroups of data. We might want to calculate means and standard deviation across genders, or countries or schools. For the audio data, we calculated the summary statistics by the word spoken:

```
group_by(word) %>%  
  summarise(m = mean(value), s = sd(value),  
            mx = max(value), mn = min(value))
```

If we want to more operations on the variable `word` after these calculations, you would need to do an `ungroup()` step.

# Counts and tallies

Dealing with categorical data, means wanting to count or tally. The function `count()` allows calculate the number of objects in levels of a variable, and the function `tally()` calculates how many objects there are. We used `count()` for the tweets example:

```
tweets %>%  
  count(source, hour = hour(with_tz(created, "EST")))
```






to calculate the number of tweets from the different devices, android or iphone, each hour.

# Putting it together for the french fries

10 week sensory experiment, 12 individuals assessed taste of french fries on several scales (how potato-y, buttery, grassy, rancid, paint-y do they taste?), fried in one of 3 different oils, replicated twice. First few rows:

```
# A tibble: 696 x 9
  time treatment subject rep potato buttery grassy rancid painty
*   <fctr>      <fctr> <fctr> <dbl>  <dbl>    <dbl>  <dbl>  <dbl>  <dbl>
1     1         1       3     1    2.9     0.0    0.0    0.0    5.5
2     1         1       3     2   14.0     0.0    0.0    1.1    0.0
3     1         1      10     1   11.0     6.4    0.0    0.0    0.0
4     1         1      10     2    9.9     5.9    2.9    2.2    0.0
5     1         1      15     1    1.2     0.1    0.0    1.1    5.1
6     1         1      15     2    8.8     3.0    3.6    1.5    2.3
7     1         1      16     1    9.0     2.6    0.4    0.1    0.2
8     1         1      16     2    8.2     4.4    0.3    1.4    4.0
9     1         1      19     1    7.0     3.2    0.0    4.9    3.2
10    1         1      19     2   13.0     0.0    3.1    4.3   10.3
# ... with 686 more rows
```

# What would we like to know?

-  Is the design complete?
-  Are replicates like each other?
-  How do the ratings on the different criteria differ?
-  Are raters giving different scores on average?
-  Do ratings change over the weeks?


Each of these questions involves different summaries of the data.

# Answer some Questions


📊 Easiest question is whether the ratings are similar on the different scales, potato'y, buttery, grassy, rancid and painty.

📊 We need to gather the data into long form, and make plots faceted by the scale.

```
ff.m <- french_fries %>%  
  gather(type, rating, -subject, -time, -treatment, -rep)  
ggplot(data=ff.m, aes(x=rating)) + geom_histogram(binwidth=2) +  
  facet_wrap(~type, ncol=5)
```




# Look at it a different way




# Comparison of the distributions of criteria

Ratings on the different criteria are quite different.

 Potato'y scores relatively highly.


 Grassy are mostly o's


 Buttery and painty are skewed right, mostly low values a few high ratings


 Rancid is quite varied, we would hope that this relates to time, that the chips get more rancid as the weeks go by.




# Do the replicates look like each other?

 We will start to tackle this by plotting the replicates against each other using a scatterplot.

 If raters give the same rating to the replicates, we would expect something close to this, then all values would lie on the  $X=Y$  line

 We need to


-  gather the data into long form, and

-  then get the replicates spread into separate columns.

```
ff.s <- ff.m %>% spread(rep, rating)
head(ff.s)
# A tibble: 6 x 6
```


	time	treatment	subject	type	`1`	`2`
	<fctr>	<fctr>	<fctr>	<chr>	<dbl>	<dbl>
1	1	1	3	buttery	0.0	0.0
2	1	1	3	grassy	0.0	0.0
3	1	1	3	painty	5.5	0.0
4	1	1	3	potato	2.9	14.0
5	1	1	3	rancid	0.0	1.1
6	1	1	10	buttery	6.4	5.9

# Check Replicates



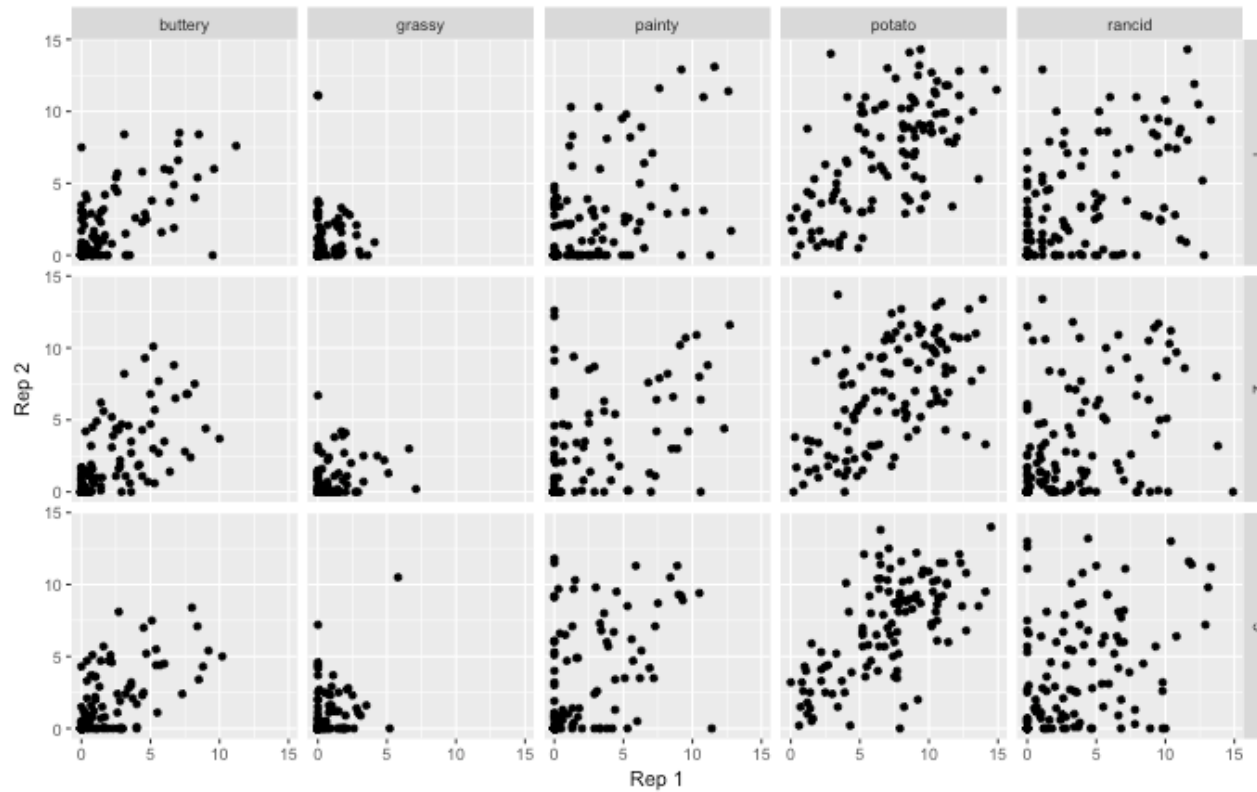
They have some positive linear association, but there is a lot more variation than expected. One rep might have scored 15 and the other 0, that's the same batches, same oil, same criteria, same week, same rater!

## Separately by criteria ...




buttery and potato'y look a bit better. The rest are still terrible.

by oil ...



Not much improvement here.

## By subject ...



Some subjects may be more experienced than others?





Because the replicates do not look like each other, the quality of the data might be questioned at this point.

Nevertheless, lets push on with some of the other questions.

# Completeness of experimental design

If the data is complete it should be  $12 \times 10 \times 3 \times 2$ , that is, 6 records for each person. (Assuming that each person rated on all scales.)

To check this we want to tabulate the number of records for each subject, time and treatment. That is,

-  select appropriate columns,
-  tabulate,
-  count and
-  spread it out to give a nice table.



```
french_fries %>%
  select(subject, time, treatment) %>%
  count(subject, time) %>%
  spread(time, n)
# A tibble: 12 x 11
  subject `1` `2` `3` `4` `5` `6` `7` `8` `9` `10`
*   <fctr> <int> <int> <int> <int> <int> <int> <int> <int> <int> <int>
1       3     6     6     6     6     6     6     6     6     6    NA
2      10     6     6     6     6     6     6     6     6     6     6
3      15     6     6     6     6     6     6     6     6     6     6
4      16     6     6     6     6     6     6     6     6     6     6
5      19     6     6     6     6     6     6     6     6     6     6
6      31     6     6     6     6     6     6     6     6     NA     6
7      51     6     6     6     6     6     6     6     6     6     6
8      52     6     6     6     6     6     6     6     6     6     6
9      63     6     6     6     6     6     6     6     6     6     6
10     78     6     6     6     6     6     6     6     6     6     6
11     79     6     6     6     6     6     6     6     6     6    NA
12     86     6     6     6     6     6     6     6     6     NA     6
```


Its pretty good, but subjects, 3, 31, 79, 86 all missed a rating session.

# By criteria

# A tibble: 12 x 11

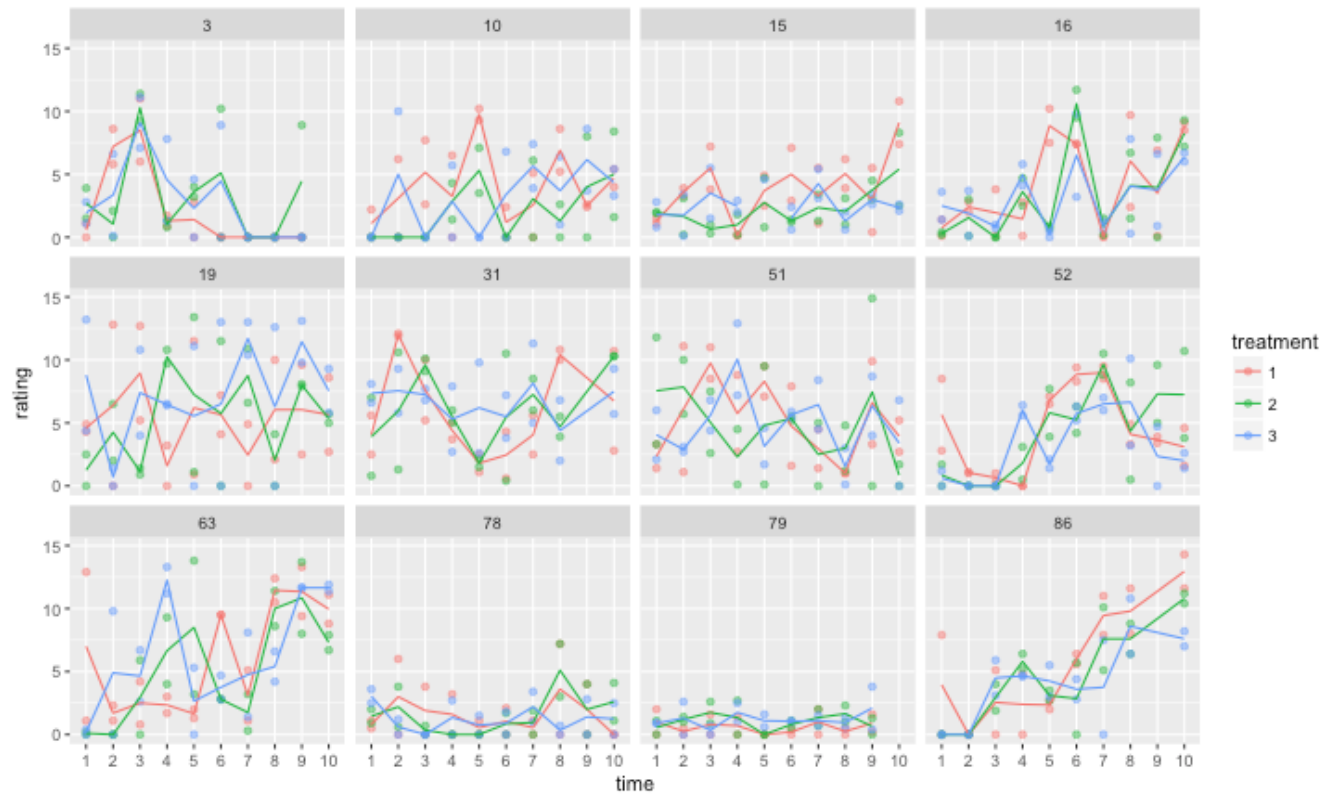
	subject	`1`	`2`	`3`	`4`	`5`	`6`	`7`	`8`	`9`	`10`
*	<fctr>	<int>	<int>	<int>	<int>	<int>	<int>	<int>	<int>	<int>	<int>
1	3	30	30	30	30	30	30	30	30	30	NA
2	10	30	30	30	30	30	30	30	30	30	30
3	15	30	30	30	30	30	30	30	30	30	30
4	16	30	30	30	30	30	30	30	30	30	30
5	19	30	30	30	30	30	30	30	30	30	30
6	31	30	30	30	30	30	30	30	30	NA	30
7	51	30	30	30	30	30	30	30	30	30	30
8	52	30	30	30	30	30	30	30	30	30	30
9	63	30	30	30	30	30	30	30	30	30	30
10	78	30	30	30	30	30	30	30	30	30	30
11	79	30	30	30	30	30	30	30	30	30	NA
12	86	30	30	30	30	30	30	30	30	NA	30

# Change in rancid ratings over weeks

 Filter on criteria

 Compute means by subject, week and oil

```
ff.av <- ff.m %>%  
  filter(type == "rancid") %>%  
  group_by(subject, time, treatment) %>%  
  summarise(rating=mean(rating))
```



Oh, its awful data! Only subject 86 is seeing the chips get more rancid over time, with maybe oil 1 being worse. Subject 63, shows some trend. Nothing is rancid for subjects 78, 79. Subject 53 thinks they taste better after many weeks in the same old oil!







# Joins

It's rare that a data analysis involves only a single table of data. Typically you have many tables of data, and you must combine them to answer the questions that you're interested in. Collectively, multiple tables of data are called relational data because it is the relations, not just the individual datasets, that are important.

```
load("data/plane_N4YRAA.rda")
plane_N4YRAA %>% glimpse()
Observations: 145
Variables: 8
$ FL_DATE    <date> 2017-05-26, 2017-05-02, 2017-05-05, 2017-05-11, 2017...
$ CARRIER   <chr> "AA", "AA", "AA", "AA", "AA", "AA", "AA", "AA", "AA",...
$ FL_NUM     <int> 2246, 2276, 2278, 2287, 2288, 2291, 2297, 2297, 2297,...
$ ORIGIN     <chr> "CVG", "DFW", "DFW", "STL", "IND", "CHS", "DFW", "DFW...
$ DEST       <chr> "DFW", "IND", "OKC", "ORD", "DFW", "DFW", "MKE", "MKE...
$ DEP_TIME   <chr> "0748", "2020", "0848", "0454", "0601", "0807", "0700...
$ ARR_TIME   <chr> "0917", "2323", "0941", "0600", "0719", "0947", "0905...
$ DISTANCE   <dbl> 812, 761, 175, 258, 761, 987, 853, 853, 853, 853, 447...
```

```
airport <- read_csv("data/airports.csv")
airport %>% select(AIRPORT, LATITUDE, LONGITUDE, AIRPORT_STATE_NAME) %>%
  glimpse()
Observations: 13,094
Variables: 4
$ AIRPORT      <chr> "01A", "03A", "04A", "05A", "06A", "07A", "...
$ LATITUDE     <dbl> 58.10944, 65.54806, 68.08333, 67.57000, 57....
$ LONGITUDE    <dbl> -152.90667, -161.07167, -163.16667, -148.18...
$ AIRPORT_STATE_NAME <chr> "Alaska", "Alaska", "Alaska", "Alaska", "Al..."
```

# Joining the two tables

-  Purpose is to show flight movement on the map
-  Key is the airport three letter code,
  -  called ORIGIN or DEST in plane\_N4YRAA table
  -  called AIRPORT in the airport table
-  One table, plane\_N4YRAA, has less airports than the other
  -  Only want to keep the rows of airport table, for those that appear in the plane\_N4YRAA table

```

airport <- airport %>%
  select(AIRPORT, LATITUDE, LONGITUDE, AIRPORT_IS_LATEST, DISPLAY_AIRPORT_NAME)
  filter(AIRPORT_IS_LATEST == 1) %>%
  select(-AIRPORT_IS_LATEST)

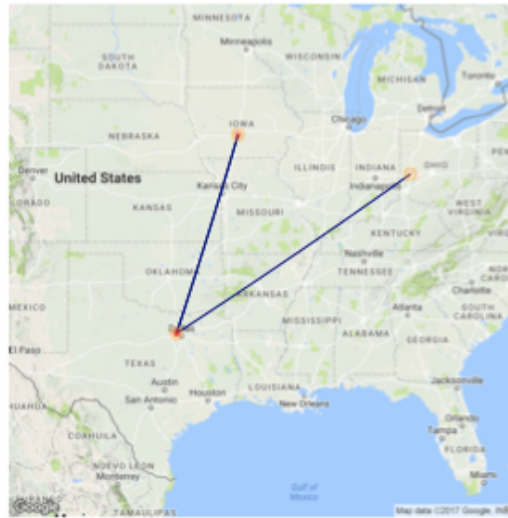
N4YRAA_latlon <- left_join(plane_N4YRAA, airport,
                           by = c("ORIGIN"="AIRPORT")) %>%
  rename("ORIGIN_LATITUDE"="LATITUDE",
         "ORIGIN_LONGITUDE"="LONGITUDE")
N4YRAA_latlon %>%
  select(ORIGIN, ORIGIN_LATITUDE, ORIGIN_LONGITUDE,
         DISPLAY_AIRPORT_NAME)
# A tibble: 146 x 4
   ORIGIN ORIGIN_LATITUDE ORIGIN_LONGITUDE
  <chr>      <dbl>          <dbl>
1   CVG      39.04889        -84.66778
2   DFW      32.89722        -97.03778
3   DFW      32.89722        -97.03778
4   STL      38.74861        -90.37000
5   IND      39.71722        -86.29472
6   CHS      32.89861        -80.04056
7   DFW      32.89722        -97.03778
8   DFW      32.89722        -97.03778
9   MKE      42.94694        -87.89694
10  MKE      42.94694        -87.89694
# ... with 136 more rows, and 1 more variables: DISPLAY_AIRPORT_NAME <chr>

```

The variables ORIGIN\_LATITUDE, ORIGIN\_LONGITUDE, DISPLAY\_AIRPORT\_NAME are added to corresponding row in the plane\_N4YRAA table.



- ▮ Added the spatial coordinates (lat, lon) for the origin airport
- ▮ The same needs to be done for the destination airport
- ▮ Then the airports can be drawn over a map



# Share and share alike



This work is licensed under a [Creative Commons Attribution 4.0 International License](#).